A PLAIN-LANGUAGE GUIDE

AI Fundamentals

Everything you need to understand modern artificial intelligence — explained without the jargon.

A friendly introduction · No technical background required

Made by Kaio to the CS Team ❤

1 What AI Actually Is

Artificial Intelligence is the broad field of getting machines to do things that normally require human intelligence — reasoning, recognizing patterns, understanding language, and making decisions.

The single most useful thing to clear up early is the nested hierarchy people constantly mix up:

- **AI** is the umbrella term for the whole field.
- **Machine Learning (ML)** is a subset of AI where the machine *learns from data* instead of following hand-written rules.
- **Deep Learning (DL)** is a subset of ML that uses **neural networks** with many layers.

> *Think of three Russian nesting dolls. AI is the big one, ML sits inside it, and DL sits inside that.*

It's also worth distinguishing **narrow AI** from **general AI (AGI)**. Everything that exists today is *narrow* — extremely good at one specific thing but useless outside that lane. A chess AI can't drive a car. AGI — a system as flexible as a human across any task — doesn't exist yet and is still theoretical. This matters because headlines blur the two and create false expectations.

2 How We Got Here

A quick history helps the whole thing make sense — AI didn't appear overnight.

- **1950s–80s — Rule-based AI.** Humans hand-wrote every rule. Powerful in narrow cases, but impossible to scale to the messy real world.
- **1990s–2000s — Machine learning takes over.** Instead of writing rules, we let computers learn from data. Spam filters and recommendations quietly became part of daily life.
- **2012 — The deep learning boom.** Neural networks crushed an image-recognition contest, kicking off the modern era. More data + more computing power made them suddenly viable.
- **2017 — The Transformer.** A new architecture (the one behind today's LLMs) made it possible to train on enormous amounts of text efficiently.
- **2022 onward — AI goes mainstream.** ChatGPT put generative AI in everyone's hands, and the tools have been getting more capable and multimodal ever since.

3 The Big Shift: From Rules to Examples

This is the conceptual heart of everything. For decades, programming meant a human writing explicit rules: *"IF the email contains 'free money' THEN mark as spam."* The problem is that the real world has millions of exceptions you can't possibly write out by hand.

Machine learning flips this. Instead of writing the rule, you **show thousands of examples** ("here are 10,000 emails, these are spam, these aren't") and the system *figures out the rule itself*.

> *You don't teach a child what a dog is by listing rules ("four legs, fur, tail"). You point at dogs over and over, and eventually they generalize. ML works the same way — it learns from exposure, not from a definition.*

4 The Three Ways Machines Learn

Supervised learning

You give labeled examples (a photo *plus* the answer "this is a cat"). The model learns to map input → correct output. This is the most common type and covers spam filters, medical image diagnosis, and price prediction. The catch: someone has to *label* all that data first, which can be slow and expensive.

Unsupervised learning

You give *unlabeled* data and let the model find structure on its own — for example, handing it customer data and letting it discover natural groupings ("these shoppers behave similarly"). Nobody told it the categories; it found them. It's how stores spot customer segments they didn't know existed.

Reinforcement learning (RL)

The model learns by *trial and error*, getting rewards for good moves and penalties for bad ones — like training a dog with treats. This is how AI mastered games like Go and chess, and a close cousin (called RLHF — reinforcement learning from human feedback) is how chatbots like ChatGPT and Claude are tuned to be helpful and polite: humans rate answers, and the model learns what "good" looks like.

> *A handy way to remember them: supervised = learning with an answer key; unsupervised = finding patterns with no answer key; reinforcement = learning from rewards and mistakes.*

5 The Building Blocks

- **Data** — the fuel. Quality and quantity matter more than almost anything else. *Garbage in, garbage out.*
- **Features** — the input variables describing each example (a house's size, location, rooms; or the pixels of an image).
- **Model** — the mathematical structure that makes predictions.
- **Parameters / weights** — the internal numbers the model adjusts while learning. This *is* the learned knowledge. Modern large models have billions of them.
- **Training** — the process of tuning those weights to reduce errors, repeated across the data.
- **Inference** — using the finished model on new data. (Training is the studying; inference is the exam.)

> *A useful frame: **data** is the textbook, **training** is the studying, the **parameters** are what got memorized, and **inference** is sitting the exam on questions never seen before.*

6 Neural Networks & Deep Learning

Neural networks are *loosely* inspired by the brain — emphasis on loosely. Picture layers of tiny decision-makers ("neurons"):

- Each connection between neurons has a **weight** (an importance value).
- A neuron combines its inputs according to those weights, runs the result through an **activation function** (which decides whether and how strongly to "fire"), and passes it forward.
- **Backpropagation** is the learning mechanism: the network compares its guess to the right answer, measures how wrong it was, and sends that error *backward* through the layers, nudging every weight toward a better answer. Repeat millions of times and the network gets good.

Why "deep"? More layers let the network build up **increasingly abstract concepts**. In image recognition, early layers detect edges, middle layers detect shapes like eyes and ears, and later layers recognize "this is a face." That stacking is the depth.

THE MOMENT IT CLICKED

Deep learning's big public breakthrough came in 2012, when a network called **AlexNet** won a major image-recognition contest by a huge margin — proving that lots of data plus powerful graphics chips (GPUs) could make these networks dramatically better. That result kicked off the modern AI era.

4

7 Generative AI & LLMs

The part everyone's talking about.

- **Generative AI** — AI that *creates* new content (text, images, code, audio) rather than just classifying. The leap from "is this spam?" to "write me an email."
- **Large Language Model (LLM)** — a model trained on enormous amounts of text whose core job is deceptively simple: **predict the next word** (technically, the next token). ChatGPT, Claude, and Gemini are all LLMs. The surprise of recent years: make this game big enough and the model develops genuinely useful abilities nobody explicitly programmed.
- **Token** — a chunk of text, roughly a word or part of one. The model thinks in tokens, not letters — which is why it can stumble on counting letters in a word.
- **Transformer** — the architecture (2017) that made modern LLMs possible. Its key innovation, the **attention mechanism**, lets the model weigh which parts of the input matter most for each word it processes.
- **Prompt** — the instruction you give the model. **Prompt engineering** is the skill of phrasing requests well.
- **Context window** — how much the model can "hold in mind" at once. Go past it and earlier parts get forgotten.
- **Hallucination** — when the model confidently produces something that *sounds* plausible but is false. An LLM is a pattern-completion engine, not a fact database. Always verify anything important.

8 Limitations & Risks

The part everyone should *talk about.*

- **Bias** — models absorb the prejudices present in their training data.
- **The black box problem** — it's often genuinely hard to explain *why* a model made a specific decision (the field of *interpretability*).
- **No real understanding** — the big one. An LLM doesn't "know" or "believe" anything; it recognizes statistical patterns in language. Mistaking fluency for understanding leads to overtrusting it.
- **Data dependence & staleness** — a model doesn't know what it wasn't trained on, and its knowledge has a cutoff date unless connected to live sources.
- **Ethical stakes** — privacy, misuse (deepfakes, scams, misinformation at scale), and economic disruption.

9 A Few More Basics Worth Knowing

Why it sometimes "makes things up"

It completes patterns, so it prefers a confident answer over admitting it doesn't know. Always check important facts.

Why it "forgets"

The context window. In a long conversation, the beginning drops off — so it helps to repeat important information.

Why the same question gives different answers

There's randomness in picking the next token (a setting called *temperature*). It isn't a database that always returns the same row.

Training vs. using

Training costs millions and takes months; using it (inference) is instant and cheap. Someone using ChatGPT isn't "training" it.

AI isn't just chatbots

Recommendations (Netflix, Spotify), spam filters, face recognition, GPS, translation. People have used AI for years without realizing it.

There's no "brain" in there

No intention, no will, no consciousness. It's math predicting patterns — and it reflects the data and human choices that shaped it, so it isn't "neutral" just because it's a machine.

10 It's Not Just Text Anymore

Early AI assistants only handled text. Today's leading models are **multimodal** — they can understand and create across several formats at once:

- **Images** — describe a photo, read a chart, or generate artwork from a description.
- **Audio** — transcribe speech, hold a spoken conversation, or generate voices and music.
- **Video** — analyze footage or generate short clips from a prompt.

So you can show a model a screenshot and ask "what's wrong here?", or hand it a PDF and ask for a summary. It's no longer just a text box.

11 Getting the Best Out of It

How you ask makes a big difference. A few habits that consistently improve results:

Give context
Say who it's for, the goal, and any constraints. "Explain to a beginner in 3 sentences" beats "explain this."

Show an example
Paste a sample of the tone or format you want. The model matches patterns well.

Break it down
Split a big task into smaller steps instead of asking for everything at once.

Ask it to think
"Walk through it step by step" often produces more accurate reasoning.

Iterate
Treat the first answer as a draft. Tell it what to change — "shorter," "more formal," "add an example."

Verify
For anything important, double-check facts, numbers, and quotes against a reliable source.

12 Privacy & Cost

MIND WHAT YOU SHARE

Assume anything you type could be stored or reviewed. Don't paste passwords, customer data, or confidential company information into AI tools unless you're using an approved, secure version with clear data policies.

And "free" isn't really free. These models run in expensive data centers, which is why there are usage limits and paid plans. Running an AI also consumes real energy and water for cooling — a growing part of the conversation around responsible use.

13 Practical Concepts & The Golden Rule

- **Fine-tuning** — taking a general pre-trained model and specializing it on a narrow domain (e.g., a medical or legal version).
- **RAG (Retrieval-Augmented Generation)** — connecting an LLM to an external knowledge base so it looks up real, current facts before answering. A major way to reduce hallucinations.
- **API** — how apps "plug into" AI without training anything themselves.
- **AI agents** — systems that use an LLM not just to *answer* but to *plan and act* across multiple steps, using tools. The frontier of where things are heading.

GOOD TASKS FOR AI

Drafting, summarizing, translating, explaining, brainstorming, and coding.

BE SKEPTICAL OF

Exact facts, numbers, dates, quotes, and anything high-stakes (health, legal, financial) without human verification.

The Golden Rule: Treat AI like a brilliant, fast intern who's sometimes overconfident — great for speeding things up, but you review before you trust.

Modern AI is extraordinarily good at recognizing and reproducing patterns, but it doesn't understand the way a person does — and knowing that difference is the key to using it well.

8